

Joining data from multiple sources

DATA 202 21FA, based on datasciencebox.org

Q&A

What does `%in%` do?

- Don't confuse with `%>%` or `==`

```
nobel %>% filter(country == c("USA", "France"))
```

(doesn't work!)

```
nobel %>% filter(country %in% c("USA", "France"))
```

```
# A tibble: 388 × 26
```

	id	firstname	surname	year	category	affiliation	city
	<dbl>	<chr>	<chr>	<dbl>	<chr>	<chr>	<chr>
1	4	Henri	Becquerel	1903	Physics	École Polytech...	Paris
2	5	Pierre	Curie	1903	Physics	École municipa...	Paris
3	6	Marie	Curie	1911	Chemistry	Sorbonne Unive...	Paris
4	11	Albert A.	Michelson	1907	Physics	University of ...	Chic...
5	12	Gabriel	Lippmann	1908	Physics	Sorbonne Unive...	Paris
6	14	Ferdinand	Braun	1909	Physics	Strasbourg Uni...	Stra...

```
# with 382 more rows and 10 more variables: country <chr>
```

Q&A

How do I remember all this syntax and keep everything straight?

It takes work. But you can do it. Tips:

- Don't "just make it work".
- Understand *why* something works (explain it!)
- Try out variations!

Data: Women in science

Information on 10 women in science who changed the world

name
Ada Lovelace
Marie Curie
Janaki Ammal
Chien-Shiung Wu
Katherine Johnson
Rosalind Franklin
Vera Rubin
Gladys West
Flossie Wong-Staal
Jennifer Doudna

Source: Discover Magazine

Inputs

professions

dates

notability

professions

```
# A tibble: 10 × 2
```

```
  name
```

```
  <chr>
```

```
profession
```

```
<chr>
```

```
1 Ada Lovelace
```

```
Mathematician
```

```
2 Marie Curie
```

```
Physicist and Chemist
```

```
3 Janaki Ammal
```

```
Botanist
```

```
4 Chien-Shiung Wu
```

```
Physicist
```

```
5 Katherine Johnson
```

```
Mathematician
```

```
6 Rosalind Franklin
```

```
Chemist
```

```
7 Vera Rubin
```

```
Astronomer
```

```
8 Gladys West
```

```
Mathematician
```

```
9 Flossie Wong-Staal
```

```
Virologist and Molecular Biologist
```

```
10 Jennifer Doudna
```

```
Biochemist
```

Desired output

```
# A tibble: 10 × 5
```

	name	profession	birth_year	death_year	known_for
	<chr>	<chr>	<int>	<int>	<chr>
1	Ada Lovelace	Mathematic...	NA	NA	first co...
2	Marie Curie	Physicist ...	NA	NA	theory o...
3	Janaki Ammal	Botanist	1897	1984	hybrid s...
4	Chien-Shiung Wu	Physicist	1912	1997	confirm a...
5	Katherine Johnson	Mathematic...	1918	2020	calculat...
6	Rosalind Franklin	Chemist	1920	1958	<NA>
7	Vera Rubin	Astronomer	1928	2016	existenc...
8	Gladys West	Mathematic...	1930	NA	mathemat...
9	Flossie Wong-Staal	Virologist...	1947	NA	first sc...
10	Jennifer Doudna	Biochemist	1964	NA	one of t...

First try: paste them together

name	profession	known_for
Ada Lovelace	Mathematician	first computer algorithm
Marie Curie	Physicist and Chemist	confirm and refine theory of radioactive beta decay, Wu experiment overturning theory of parity
Janaki Ammal	Botanist	first scientist to clone HIV and create a map of its genes which led to a test for the virus
Chien-Shiung Wu	Physicist	mathematical modeling of the shape of the Earth which served as the foundation of GPS technology
Katherine Johnson	Mathematician	hybrid species, biodiversity protection
Rosalind Franklin	Chemist	one of the primary developers of CRISPR, a ground-breaking technology for editing genomes
Vera Rubin	Astronomer	calculations of orbital mechanics critical to sending the first Americans into space
Gladys	Mathematician	theory of radioactivity, discovery of elements polonium and radium, first woman to win a Nobel

What was wrong?

What was wrong?

How do we know which rows match up?

What was wrong?

How do we know which rows match up?

Need a **key**.

What can serve as a **key** for this data?

Mutating joins

- I have a data frame `x`
- I want extra information about things in `x`
- Some other table, `y`, has that information.
- A "join" lets me look it up.

x		y	
1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

Graphics thanks to [tidyexplain](#)

Types of joins

If `x` and `y` match up one-to-one, no difference.

What to do when things don't exactly line up?

- **full** or **outer** join: Leave blanks (`NA`) for mismatches
- **inner** join: Drop rows with any mismatches
- **left** / **right** join: Drop rows where one of the sides has a mismatch

Setup

x

```
# A tibble: 3 × 2
  key xdata
  <dbl> <chr>
1     1 x1
2     2 x2
3     3 x3
```

y

```
# A tibble: 3 × 2
  key ydata
  <dbl> <chr>
1     1 y1
2     2 y2
3     4 y4
```

x

1	x1
2	x2
3	x3

y

1	y1
2	y2
4	y4

left_join()

All rows from *x*.

```
left_join(x, y, by = "key")
```

```
# A tibble: 3 × 3  
  key xdata ydata  
  <dbl> <chr> <chr>  
1     1 x1    y1  
2     2 x2    y2  
3     3 x3    <NA>
```

left_join(x, y)

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

left_join()

```
professions %>%  
  left_join(dates)
```

```
# A tibble: 10 × 4
```

	name <chr>	profession <chr>	birth_year <int>	death_year <int>
1	Ada Lovelace	Mathematician	NA	NA
2	Marie Curie	Physicist and Chemist	NA	NA
3	Janaki Ammal	Botanist	1897	1984
4	Chien-Shiung Wu	Physicist	1912	1997
5	Katherine Johnson	Mathematician	1918	2020
6	Rosalind Franklin	Chemist	1920	1958
7	Vera Rubin	Astronomer	1928	2016
8	Gladys West	Mathematician	1930	NA
9	Flossie Wong-Staal	Virologist and Molec...	1947	NA
10	Jennifer Doudna	Biochemist	1964	NA

right_join()

All rows from *y*.

```
right_join(x, y)
```

```
# A tibble: 3 × 3  
  key xdata ydata  
  <dbl> <chr> <chr>  
1     1 x1    y1  
2     2 x2    y2  
3     4 <NA> y4
```

right_join(x, y)

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

right_join()

```
professions %>%  
  right_join(dates)
```

```
# A tibble: 8 × 4
```

	name <chr>	profession <chr>	birth_year <int>	death_year <int>
1	Janaki Ammal	Botanist	1897	1984
2	Chien-Shiung Wu	Physicist	1912	1997
3	Katherine Johnson	Mathematician	1918	2020
4	Rosalind Franklin	Chemist	1920	1958
5	Vera Rubin	Astronomer	1928	2016
6	Gladys West	Mathematician	1930	NA
7	Flossie Wong-Staal	Virologist and Molecu...	1947	NA
8	Jennifer Doudna	Biochemist	1964	NA

full_join()

All rows from both `x` and `y`. Leave `NA` for mismatches.

```
full_join(x, y, by = "key")
```

```
# A tibble: 4 × 3  
  key xdata ydata  
  <dbl> <chr> <chr>  
1     1 x1    y1  
2     2 x2    y2  
3     3 x3    <NA>  
4     4 <NA> y4
```

full_join(x, y)

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

full_join()

```
dates %>%  
  full_join(notability)
```

```
# A tibble: 10 × 4
```

	name <chr>	birth_year <int>	death_year <int>	known_for <chr>
1	Janaki Ammal	1897	1984	hybrid species, biod...
2	Chien-Shiung Wu	1912	1997	confim and refine th...
3	Katherine Johnson	1918	2020	calculations of orbi...
4	Rosalind Franklin	1920	1958	<NA>
5	Vera Rubin	1928	2016	existence of dark ma...
6	Gladys West	1930	NA	mathematical modelin...
7	Flossie Wong-Staal	1947	NA	first scientist to c...
8	Jennifer Doudna	1964	NA	one of the primary d...
9	Ada Lovelace	NA	NA	first computer algor...
10	Marie Curie	NA	NA	theory of radioactiv...

inner_join()

All matching rows. Drops mismatches.

```
inner_join(x, y, by = "key")
```

```
# A tibble: 2 × 3  
  key xdata ydata  
  <dbl> <chr> <chr>  
1     1 x1    y1  
2     2 x2    y2
```

inner_join(x, y)

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

inner_join()

```
dates %>%
```

```
  inner_join(notability)
```

```
# A tibble: 7 × 4
```

	name	birth_year	death_year	known_for
	<chr>	<int>	<int>	<chr>
1	Janaki Ammal	1897	1984	hybrid species, biodi...
2	Chien-Shiung Wu	1912	1997	confim and refine the...
3	Katherine Johnson	1918	2020	calculations of orbit...
4	Vera Rubin	1928	2016	existence of dark mat...
5	Gladys West	1930	NA	mathematical modeling...
6	Flossie Wong-Staal	1947	NA	first scientist to cl...
7	Jennifer Doudna	1964	NA	one of the primary de...

Summary of Mutating Joins

- `full_join()`: all rows from both `x` and `y`
- `inner_join()`: all *matching* rows from `x` where there are matching values in `y`.
- `left_join()`: all rows from `x`
- `right_join()`: all rows from `y`

Multiple matches? Return all combinations.

We want to get the dates and works of all the scientists. Which join function should we use?

```
# A tibble: 10 × 5
```

	name <chr>	profession <chr>	birth_year <int>	death_year <int>	known_for <chr>
1	Ada Lovelace	Mathematic...	NA	NA	first co...
2	Marie Curie	Physicist ...	NA	NA	theory o...
3	Janaki Ammal	Botanist	1897	1984	hybrid s...
4	Chien-Shiung Wu	Physicist	1912	1997	confim a...
5	Katherine Johnson	Mathematic...	1918	2020	calculat...
6	Rosalind Franklin	Chemist	1920	1958	<NA>
7	Vera Rubin	Astronomer	1928	2016	existenc...
8	Gladys West	Mathematic...	1930	NA	mathemat...
9	Flossie Wong-Staal	Virologist...	1947	NA	first sc...
10	Jennifer Doudna	Biochemist	1964	NA	one of t...

```
names(professions)
```

```
[1] "name"      "profession"
```

```
nrow(professions)
```

```
[1] 10
```

```
names(dates)
```

```
[1] "name"      "birth_year" "
```

```
nrow(dates)
```

```
[1] 8
```

```
names(notability)
```

```
[1] "name"      "known_for"
```

```
nrow(notability)
```

```
[1] 9
```



```
professions %>%
  left_join(dates) %>%
  left_join(notability)
```

```
# A tibble: 10 × 5
```

	name	profession	birth_year	death_year	known_for
	<chr>	<chr>	<int>	<int>	<chr>
1	Ada Lovelace	Mathematic...	NA	NA	first co...
2	Marie Curie	Physicist ...	NA	NA	theory o...
3	Janaki Ammal	Botanist	1897	1984	hybrid s...
4	Chien-Shiung Wu	Physicist	1912	1997	confim a...
5	Katherine Johnson	Mathematic...	1918	2020	calculat...
6	Rosalind Franklin	Chemist	1920	1958	<NA>
7	Vera Rubin	Astronomer	1928	2016	existenc...
8	Gladys West	Mathematic...	1930	NA	mathemat...
9	Flossie Wong-Staal	Virologist...	1947	NA	first sc...
10	Jennifer Doudna	Biochemist	1964	NA	one of t...

Case study: Grocery sales

Grocery sales

- Have:
 - *Purchases*: One row per customer per item, listing purchases they made
 - *Prices*: One row per item in the store, listing their prices
- **Want**: Total revenue

purchases

customer_id	item
c1	bread
c1	milk
c1	banana
c2	milk
c2	toilet paper

prices

item	price
avocado	0.50
banana	0.15
bread	1.00
milk	0.80
toilet paper	3.00

Grocery sales

```
purchases %>%  
  left_join(prices)
```

```
# A tibble: 5 × 3  
  customer_id item      price  
  <chr>      <chr>    <dbl>  
1 c1        bread      1  
2 c1        milk      0.8  
3 c1        banana    0.15  
4 c2        milk      0.8  
5 c2        toilet paper 3
```

Grocery sales

Total revenue

Revenue per customer

```
purchases %>%  
  left_join(prices)
```

```
# A tibble: 5 × 3  
  customer_id item      price  
  <chr>      <chr>    <dbl>  
1 c1        bread      1  
2 c1        milk      0.8  
3 c1        banana    0.1  
4 c2        milk      0.8  
5 c2        toilet paper 3
```

```
purchases %>%  
  left_join(prices) %>%  
  summarize(total_revenue = sum(price))
```

```
# A tibble: 1 × 1  
  total_revenue  
  <dbl>  
1          5.75
```

Extension: Multiple matching rows

professions_multi

dates

notability_multi

professions_multi

```
# A tibble: 12 × 2
  name           profession
<chr>          <chr>
1 Ada Lovelace   Mathematician
2 Marie Curie    Physicist
3 Marie Curie    Chemist
4 Janaki Ammal   Botanist
5 Chien-Shiung Wu Physicist
6 Katherine Johnson Mathematician
# ... with 6 more rows
```

Other types of joins

These are less common:

filtering joins

- `semi_join()`: include a row from `x` only if there's some match in `y`
- `anti_join()`: include a row from `x` only if there's *no* match in `y`

nest join

- `nest_join()`: get bundles of all matching rows from `y` (most flexible)

Specifying keys

- Keys must match *exactly*
- Can join on multiple columns (first name **and** last name)
- Default join: columns with same names
- Specify what columns to use: `left_join(x, y, by = c("first_name", "last_name"))`

General notes

A note on mutate

- Badly named. Think "add_computed_column".
- **DON'T** think of it operating on a variable ("mutate the ride's start time").
- **DO** think of it operating on a data frame ("add a column computed by flooring the start time")

Speaking of better names

| `select` vs `filter`?

Maybe should have been named `select_columns` and `select_rows`.